

# Instacart Market Basket Analysis EDA

*Xi Liang*

*6/26/2017*

## Contents

<b>Loading files</b>	<b>2</b>
<b>Departments</b>	<b>2</b>
<b>Aisles</b>	<b>3</b>
<b>Products</b>	<b>4</b>
Products with missing information . . . . .	5
<b>Orders</b>	<b>10</b>
NA values . . . . .	11
Number of transaction made . . . . .	12
Order day of week and hour of the day . . . . .	13
Order Hour of the Day . . . . .	14
Combining Day of Week and Hour of the Day . . . . .	15
Combining “Days_since_prior_order” and Transaction Made . . . . .	16
Distribution of Days Since Prior Order . . . . .	17
When customers come back, do they make purchase around the same time (day/hour) ? . . . . .	18
<b>Orders_products (prior and train)</b>	<b>20</b>
Most popular products . . . . .	21
First product add to the cart . . . . .	21
Which products get reordered most . . . . .	22
<b>In-depth Analysis</b>	<b>22</b>
Combining everything . . . . .	22
Most popular products in each department . . . . .	24

```
library(data.table)
```

```
library(dplyr)
```

```
## -----
```

```
## data.table + dplyr code now lives in dtplyr.
```

```
## Please library(dtplyr)!
```

```
## -----
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:data.table':
```

```
##
```

```
##     between, first, last
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##     filter, lag
```

```
## The following objects are masked from 'package:base':
##
## intersect, setdiff, setequal, union

library(tm)

## Loading required package: NLP

library(SnowballC)
library(wordcloud)

## Loading required package: RColorBrewer

library(ggplot2)

##
## Attaching package: 'ggplot2'

## The following object is masked from 'package:NLP':
##
## annotate

library(treemap)
```

## Loading files

```
aisles <- fread("data/aisles.csv", stringsAsFactors = T)
dept <- fread("data/departments.csv", stringsAsFactors = T)
orders <- fread("data/orders.csv", stringsAsFactors = T)
products <- fread("data/products.csv", stringsAsFactors = T)
order_products_prior <- fread("data/order_products__prior.csv", stringsAsFactors = T)

##
Read 23.8% of 32434489 rows
Read 54.6% of 32434489 rows
Read 83.7% of 32434489 rows
Read 32434489 rows and 4 (of 4) columns from 0.538 GB file in 00:00:05

order_products_train <- fread("data/order_products__train.csv", stringsAsFactors = T)
```

This dataset contains total six files, we will take a look at them one by one.

## Departments

```
dept$department

## [1] frozen      other        bakery       produce
## [5] alcohol     international beverages     pets
## [9] dry goods  pasta bulk      personal care meat seafood
## [13] pantry      breakfast    canned goods dairy eggs
## [17] household   babies       snacks       deli
## [21] missing
## 21 Levels: alcohol babies bakery beverages breakfast bulk ... snacks
```

“Dept” data contains the names for 20 different departments, mostly are department names we see in our day to day life in a grocery store. This file also contain a category, “missing”, to describe items that are not associated to any departments listed above.

## Aisles

```
aisles$aisle
```

## [1] prepared soups salads	specialty cheeses
## [3] energy granola bars	instant foods
## [5] marinades meat preparation	other
## [7] packaged meat	bakery desserts
## [9] pasta sauce	kitchen supplies
## [11] cold flu allergy	fresh pasta
## [13] prepared meals	tofu meat alternatives
## [15] packaged seafood	fresh herbs
## [17] baking ingredients	bulk dried fruits vegetables
## [19] oils vinegars	oral hygiene
## [21] packaged cheese	hair care
## [23] popcorn jerky	fresh fruits
## [25] soap	coffee
## [27] beers coolers	red wines
## [29] honeys syrups nectars	latino foods
## [31] refrigerated	packaged produce
## [33] kosher foods	frozen meat seafood
## [35] poultry counter	butter
## [37] ice cream ice	frozen meals
## [39] seafood counter	dog food care
## [41] cat food care	frozen vegan vegetarian
## [43] buns rolls	eye ear care
## [45] candy chocolate	mint gum
## [47] vitamins supplements	breakfast bars pastries
## [49] packaged poultry	fruit vegetable snacks
## [51] preserved dips spreads	frozen breakfast
## [53] cream	paper goods
## [55] shave needs	diapers wipes
## [57] granola	frozen breads doughs
## [59] canned meals beans	trash bags liners
## [61] cookies cakes	white wines
## [63] grains rice dried goods	energy sports drinks
## [65] protein meal replacements	asian foods
## [67] fresh dips tapenades	bulk grains rice dried goods
## [69] soup broth bouillon	digestion
## [71] refrigerated pudding desserts	condiments
## [73] facial care	dish detergents
## [75] laundry	indian foods
## [77] soft drinks	crackers
## [79] frozen pizza	deodorants
## [81] canned jarred vegetables	baby accessories
## [83] fresh vegetables	milk
## [85] food storage	eggs
## [87] more household	spreads

```
## [89] salad dressing toppings      cocoa drink mixes
## [91] soy lactosefree              baby food formula
## [93] breakfast bakery             tea
## [95] canned meat seafood          lunch meat
## [97] baking supplies decor        juice nectars
## [99] canned fruit applesauce      missing
## [101] air fresheners candles       baby bath body care
## [103] ice cream toppings           spices seasonings
## [105] doughs gelatins bake mixes   hot dogs bacon sausage
## [107] chips pretzels               other creams cheeses
## [109] skin care                    pickled goods olives
## [111] plates bowls cups flatware   bread
## [113] frozen juice                  cleaning products
## [115] water seltzer sparkling water frozen produce
## [117] nuts seeds dried fruit       first aid
## [119] frozen dessert               yogurt
## [121] cereal                       meat counter
## [123] packaged vegetables fruits   spirits
## [125] trail mix snack mix          feminine care
## [127] body lotions soap           tortillas flat bread
## [129] frozen appetizers sides      hot cereal pancake mixes
## [131] dry pasta                     beauty
## [133] muscles joints pain relief    specialty wines champagnes
## 134 Levels: air fresheners candles asian foods ... yogurt
```

This data contains 134 observations that describe the locations of a specific product in the market.

## Products

```
products %>% glimpse()
```

```
## Observations: 49,688
## Variables: 4
## $ product_id    <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 1...
## $ product_name  <fctr> Chocolate Sandwich Cookies, All-Seasons Salt, R...
## $ aisle_id      <int> 61, 104, 94, 38, 5, 11, 98, 116, 120, 115, 31, 1...
## $ department_id <int> 19, 13, 7, 1, 13, 11, 7, 1, 16, 7, 7, 1, 11, 17,...
```

“Products” contains 49688 unique observations (products) and each of the product is linked with specific aisle ID and department ID. Since we already have the files that contain department and aisle information, we could consolidate those information into this file to gain a better understanding the inventory of our grocery.

```
products_w_desc <- products
```

```
products_w_desc[, aisles_description := aisles$aisle[products$aisle_id]]
products_w_desc[, depart_description := dept$department[products$department_id]]
```

Here is what our data frame looks like after integrating the aisle and department information:

```
products_w_desc
```

```
##      product_id
## 1:             1
## 2:             2
## 3:             3
```

```

##      4:      4
##      5:      5
##    ---
## 49684:    49684
## 49685:    49685
## 49686:    49686
## 49687:    49687
## 49688:    49688
##
##                                product_name
##      1:    Chocolate Sandwich Cookies
##      2:                                All-Seasons Salt
##      3:                                Robust Golden Unsweetened Oolong Tea
##      4: Smart Ones Classic Favorites Mini Rigatoni With Vodka Cream Sauce
##      5:                                Green Chile Anytime Sauce
##    ---
## 49684:                                Vodka, Triple Distilled, Twist of Vanilla
## 49685:                                En Crouete Roast Hazelnut Cranberry
## 49686:                                Artisan Baguette
## 49687:    Smartblend Healthy Metabolism Dry Cat Food
## 49688:                                Fresh Foaming Cleanser
##      aisle_id department_id      aisles_description
##      1:      61          19      cookies cakes
##      2:     104          13      spices seasonings
##      3:      94           7              tea
##      4:      38           1      frozen meals
##      5:       5          13 marinades meat preparation
##    ---
## 49684:     124           5              spirits
## 49685:      42           1    frozen vegan vegetarian
## 49686:     112           3              bread
## 49687:      41           8      cat food care
## 49688:      73          11      facial care
##      depart_description
##      1:      snacks
##      2:      pantry
##      3:    beverages
##      4:      frozen
##      5:      pantry
##    ---
## 49684:      alcohol
## 49685:      frozen
## 49686:      bakery
## 49687:      pets
## 49688:    personal care

```

## Products with missing information

Let's first take a look how many products in our data that don't have aisle and department information, and we would like to know what those products are. This will be a bit similar to text mining, so let's handle it with wordcloud.

```

products_missing_info <- products_w_desc[aisles_description == "missing" | depart_description == "missing"]
products_missing_info %>% dim

```

```
## [1] 1258      6
products_corpus <- VCorpus(VectorSource(products_missing_info$product_name))

#cleaning text info
products_corpus_clean <- tm_map(products_corpus,
                                content_transformer(tolower))

#check if it worked
products_corpus_clean[[1]] %>% as.character

## [1] "ultra antibacterial dish liquid"

#remove numbers from corpus
products_corpus_clean <- tm_map(products_corpus_clean, removeNumbers)


#remove stop words (assuming if there are any)
products_corpus_clean <- tm_map(products_corpus_clean, removeWords, stopwords())

#remove punctuation
products_corpus_clean <- tm_map(products_corpus_clean, removePunctuation)

#remove white spaces
products_corpus_clean <- tm_map(products_corpus_clean, stripWhitespace)
```

We will visualize the most frequent appeared words (at least appeared 25 times in the data) with wordcloud.

```
wordcloud(products_corpus_clean, min.freq = 25, random.order = F)
```




```
products_corpus_clean_dtm <- TermDocumentMatrix(products_corpus_clean)
mat <- as.matrix(products_corpus_clean_dtm)
v <- sort(rowSums(mat), decreasing = TRUE)
d <- data.frame(word = names(v), freq = v)
head(d,20)
```

```
##           word freq
## organic      organic 175
## yogurt       yogurt  92
## chocolate    chocolate 82
## cheese       cheese  59
```

```
## chicken      chicken  45
## vanilla      vanilla  45
## milk         milk    43
## strawberry   strawberry 41
## apple        apple   38
## coffee       coffee   35
## cream        cream   35
## greek        greek    35
## natural      natural  33
## free         free     31
## almond       almond   30
## coconut      coconut  30
## original     original  30
## potato       potato   30
## fruit        fruit    29
## white        white    28
```

```
barplot(d[1:15,]$freq, las = 2 ,names.arg = d[1:15,]$word,
        main = "Products with Missing Information, In Decreasing Frequency",
        ylab = "Frequency")
```

### Products with Missing Information, In Decreasing Frequency



After analyzing what are the products that have highest rate of missing information, we will analyze the rest of the products in the grocery that have both of the aisle and department information. Through this analysis, we will know which departments contain the most products and what are those products.

```
count_by_depart <- table(products_w_desc$depart_description) %>%
  sort(decreasing = T) %>%
  as.data.frame()

colnames(count_by_depart) <- c("department_name", "Count")
```

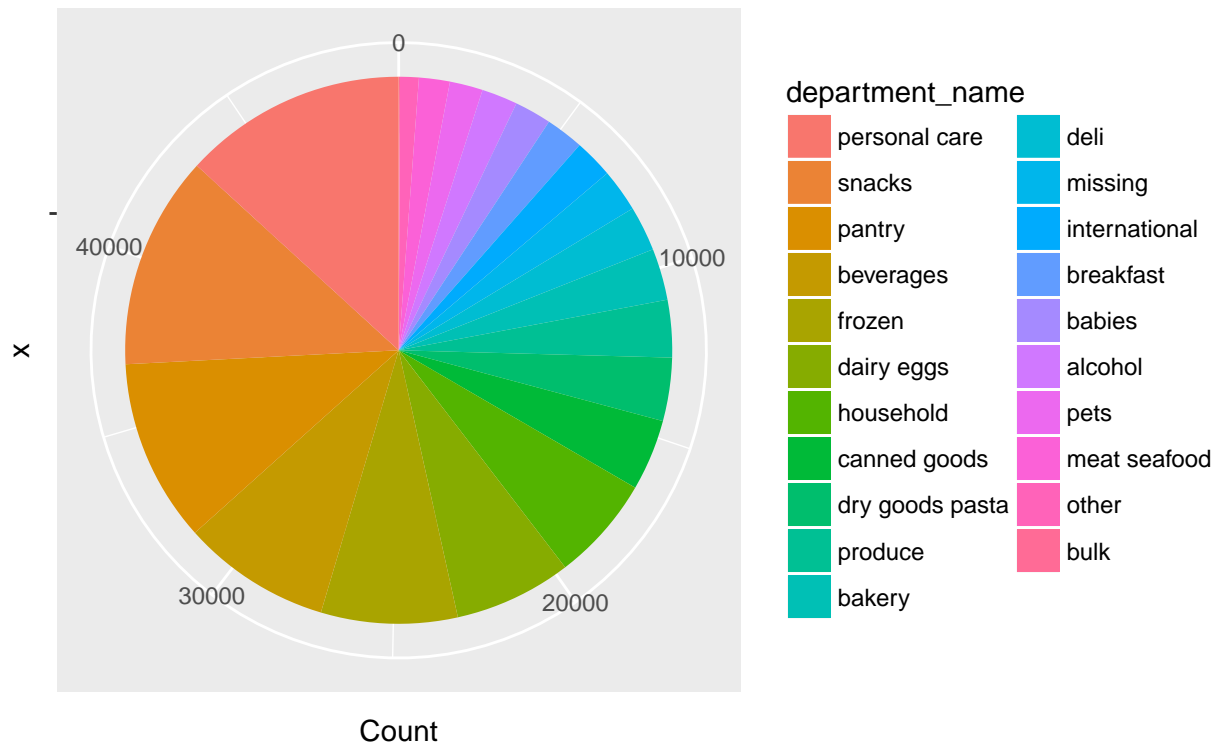
```
count_by_depart$pct <- prop.table(count_by_depart$Count) %>%
  round(3) * 100
```

```
count_by_depart
```

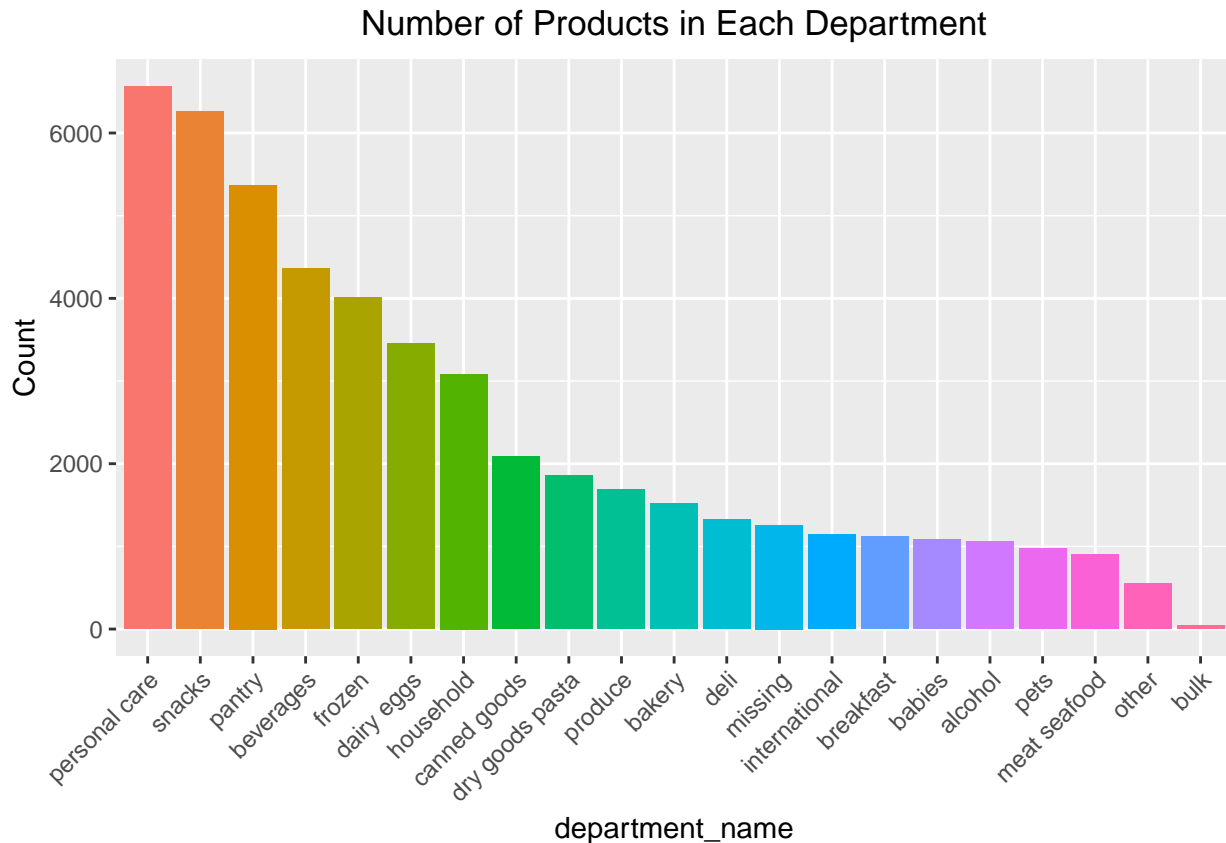
```
##   department_name Count  pct
## 1   personal care  6563 13.2
## 2      snacks    6264 12.6
## 3     pantry    5371 10.8
## 4   beverages   4365  8.8
## 5     frozen   4007  8.1
## 6   dairy eggs  3449  6.9
## 7   household  3085  6.2
## 8   canned goods 2092  4.2
## 9 dry goods pasta 1858  3.7
## 10    produce   1684  3.4
## 11     bakery   1516  3.1
## 12      deli    1322  2.7
## 13    missing   1258  2.5
## 14 international 1139  2.3
## 15    breakfast 1115  2.2
## 16     babies   1081  2.2
## 17    alcohol   1054  2.1
## 18      pets    972  2.0
## 19  meat seafood  907  1.8
## 20      other   548  1.1
## 21      bulk    38  0.1
```

```
par(mfrow = c(1,2))
ggplot(count_by_depart, aes(x = "", y = Count, fill = department_name)) +
  geom_bar(width = 1, stat = "identity") +
  coord_polar("y", start = 0)
```





```
ggplot(count_by_depart, aes(department_name, Count)) +
  geom_bar(aes(fill = department_name), stat = "identity") +
  labs(title = "Number of Products in Each Department") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1),
        plot.title = element_text(hjust = 0.5),
        legend.position = "none")
```



From the above result, we see department that has the most product is the personal care department (>6000 products), following by snacks. We will make a wild assumption by assuming that the number of product that a department carries should be correlated to the sales. For now, we will move on to the next data.

## Orders

This data contains detail purchase history of customers. There are approximately 3.4 million transaction history and 206,209 customers. Variables include “order\_id”, “user\_id”, “eval\_set”, “order\_number”, “order\_dow”, “order\_hour\_of\_day”, and “days\_since\_prior\_order”.

```
orders %>% dim
```

```
## [1] 3421083      7
```

```
unique(orders$user_id) %>% length
```

```
## [1] 206209
```

```
str(orders)
```

```
## Classes 'data.table' and 'data.frame': 3421083 obs. of 7 variables:
```

```
## $ order_id      : int  2539329 2398795 473747 2254736 431534 3367565 550135 3108588 2295261
## $ user_id       : int   1 1 1 1 1 1 1 1 1 1 ...
## $ eval_set      : Factor w/ 3 levels "prior","test",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ order_number  : int   1 2 3 4 5 6 7 8 9 10 ...
## $ order_dow     : int   2 3 3 4 4 2 1 1 1 4 ...
## $ order_hour_of_day : int   8 7 12 7 15 7 9 14 16 8 ...
## $ days_since_prior_order: num  NA 15 21 29 28 19 20 14 0 30 ...
```

```
## - attr(*, ".internal.selfref")=<externalptr>
```

Before doing anything, we want to set a few variables in the data as factors.

```
orders$order_id <- as.factor(orders$order_id)
orders$user_id <- as.factor(orders$user_id)
orders$order_hour_of_day <- as.factor(orders$order_hour_of_day)
```

Among all the variables, we see a few interesting things. From the summary, we know that all the order IDs are unique; the maximum transaction history that one has is 100; “days\_since\_prior\_order” has almost 200K records are NAs.

```
summary(orders)
```

```
##      order_id      user_id      eval_set      order_number
## 1      :      1 210      :      100 prior:3214874 Min.   : 1.00
## 2      :      1 310      :      100 test  : 75000 1st Qu.: 5.00
## 3      :      1 313      :      100 train: 131209 Median : 11.00
## 4      :      1 690      :      100          Mean  : 17.15
## 5      :      1 786      :      100          3rd Qu.: 23.00
## 6      :      1 964      :      100          Max.   :100.00
## (Other):3421077 (Other):3420483
##      order_dow      order_hour_of_day days_since_prior_order
## Min.   :0.000 10      : 288418 Min.   : 0.00
## 1st Qu.:1.000 11      : 284728 1st Qu.: 4.00
## Median :3.000 15      : 283639 Median : 7.00
## Mean   :2.776 14      : 283042 Mean   :11.11
## 3rd Qu.:5.000 13      : 277999 3rd Qu.:15.00
## Max.   :6.000 12      : 272841 Max.   :30.00
##          (Other):1730416 NA's   :206209
```

## NA values

Let's find out what those missing values have in common.

```
missing_index <- which(is.na(orders$days_since_prior_order))
missing_df <- orders[missing_index,] %>% tbl_df
missing_df %>% summary
```

```
##      order_id      user_id      eval_set      order_number
## 20      :      1 1      :      1 prior:206209 Min.   :1
## 35      :      1 2      :      1 test  : 0 1st Qu.:1
## 37      :      1 3      :      1 train: 0 Median :1
## 57      :      1 4      :      1          Mean  :1
## 75      :      1 5      :      1          3rd Qu.:1
## 100     :      1 6      :      1          Max.   :1
## (Other):206203 (Other):206203
##      order_dow      order_hour_of_day days_since_prior_order
## Min.   :0.000 15      : 17264 Min.   : NA
## 1st Qu.:1.000 16      : 17094 1st Qu.: NA
## Median :3.000 14      : 17035 Median : NA
## Mean   :2.754 12      : 16995 Mean   :NaN
## 3rd Qu.:5.000 11      : 16916 3rd Qu.: NA
## Max.   :6.000 13      : 16912 Max.   : NA
##          (Other):103993 NA's   :206209
```

As we see from the summary above, all these missing values come from order number is one. Which make sense, as there will not be transaction history of a new customer. Next, let us see the transaction quantity and its distribution.

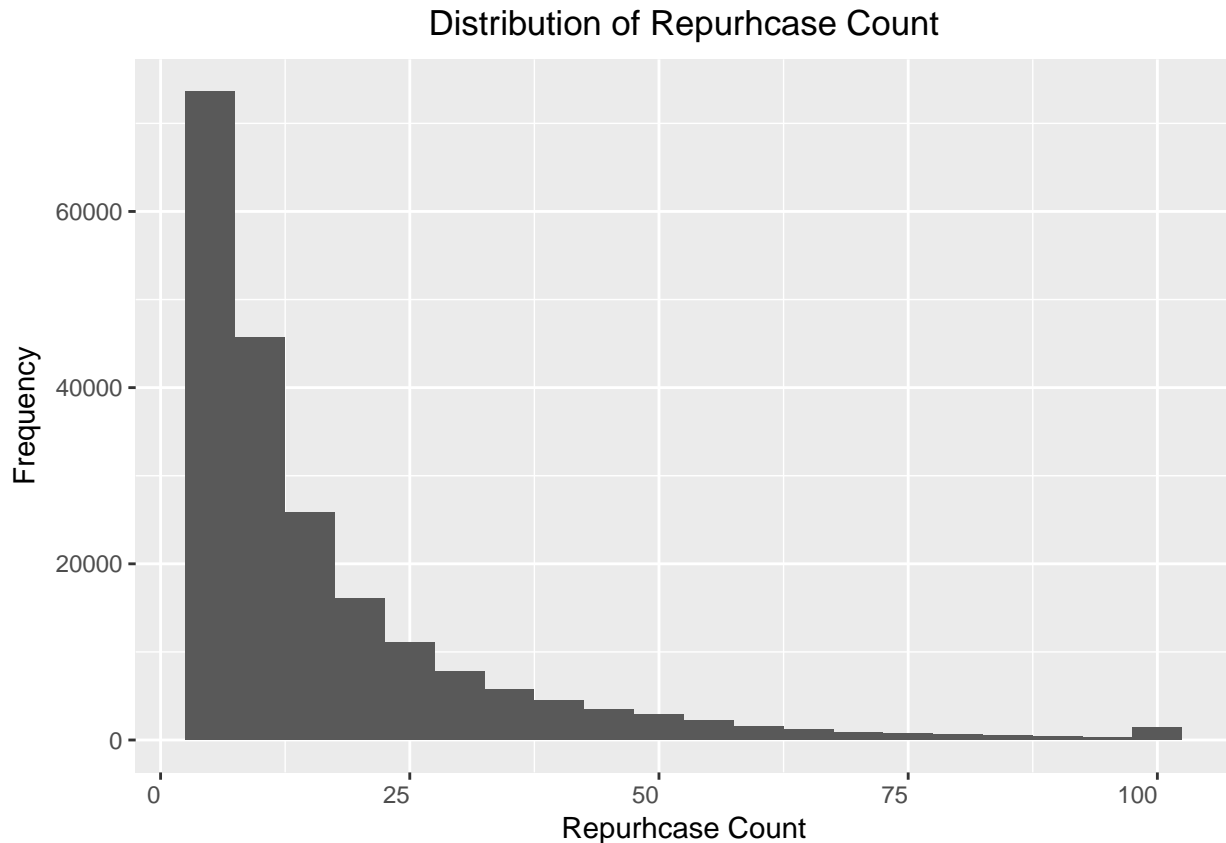
## Number of transaction made

```
transaction_count <- orders[, .(transaction_count = .N), by = user_id]
transaction_count %>% summary
```

```
##      user_id      transaction_count
## 1      :      1   Min.      : 4.00
## 2      :      1   1st Qu.: 6.00
## 3      :      1   Median : 10.00
## 4      :      1   Mean    : 16.59
## 5      :      1   3rd Qu.: 20.00
## 6      :      1   Max.     :100.00
## (Other):206203
```

We see that the minimum transaction made by specific customers in this data is 4, and the maximum is 100, with mean equals to 16.59 and median equals to 10. Following is the distribution of the transaction count frequency:

```
ggplot(transaction_count, aes(transaction_count)) +
  geom_histogram(binwidth = 5) +
  labs(title = "Distribution of Repurchase Count") +
  xlab("Repurchase Count") +
  ylab("Frequency") +
  theme(axis.text.x = element_text(hjust = 1),
        plot.title = element_text(hjust = 0.5))
```



A large portion of customer made less than 25 transactions, and that number continue to decrease as the transaction count increases. The population of customers that made 4 to 5 purchases from the store is the largest.

```
transaction_count[transaction_count == 100] %>% nrow
```

```
## [1] 1374
```


We have 1374 customers that made 100 transactions. It will be interesting to see what kind of products they bought and try to learn thier purchasing behaviors.

## Order day of week and hour of the day

Our data also includes what day of the week and what time of the day that a specific transaction happened. Let's take a look how these number fluctuate throughout the week.

```
dow <- orders[, .N, by = order_dow][order(order_dow)]
plot(dow, type = "l",
     main = "Order Count Based on Day of Week",
     xlab = "Day of the Week",
     ylab = "Count")
```


## Order Count Based on Day of Week



From the plot above, we can see that day 0 and day 1 of the week are the peak days where the customers made most of their purchases, and the count decreases as the week develops, and eventually bounce back on day 5.

## Order Hour of the Day

```
hod <- orders[, .N, by = order_hour_of_day][order(order_hour_of_day)]
ggplot(hod, aes(x= order_hour_of_day, y= N, group = 1)) +
  geom_point() +
  geom_line() +
  xlab("Hour of Day") +
  ylab("Count") +
  ggtitle("Order Count Based on Hour of Day") +
  theme(plot.title = element_text(hjust = 0.5))
```




The peak hours are from 10AM to 4PM everyday. Since we don't have the locations where the customers are, so we wouldn't know if there is a difference in peak hours across the nation.

### Combining Day of Week and Hour of the Day

```
order_dow_hod <- orders[, .N, by = .(order_dow, order_hour_of_day)]

ggplot(order_dow_hod, aes(as.numeric(order_hour_of_day), N, color = factor(order_dow))) +
  geom_line() +
  labs(title = "Transaction Count Throughout the Day, Colored by Day of the Week") +
  theme(plot.title = element_text(hjust = 0.5)) +
  xlab("Hour of Day") +
  ylab("Count")
```

## Transaction Count Throughout the Day, Colored by Day of the Week



Combining all days in one week, we can see how they are different and similar to each other. As we can see, the pattern from 1AM to 7AM are pretty similar no matter which day it is throughout the week. After 7AM, both day 0 and day 1 start to pick up speed and eventually reach approximately 600K transactions at 3PM (day 0) and 11AM (day 1). We can also see that day 0 and day 5 have similar pattern at peak hours, despite the vast difference in transaction counts, once the transaction quantity go up after 7AM, they don't drop down until it reach the daily maximum. Comparing to the other days (all days beside day 1), we see these bumps created by decline in sales volume after 11AM, and start to bounce back up after 1PM, and eventually died down around 5PM.

## Combining “Days\_since\_prior\_order” and Transaction Made

Since we can get the number of transaction of a specific customer made based on the number of occurrence of a specific “user\_id”, and because that variable “days\_since\_prior\_order” indicates the purchase frequency of a customer, it will be interesting to look at the relationship between these two variables.

```
count_freq <- transaction_count[order(user_id)]


freq_df <- orders[, .(freq = median(days_since_prior_order, na.rm = T)), by = user_id]

count_freq$freq <- freq_df$freq

ggplot(count_freq, aes(transaction_count, freq)) +
  geom_smooth() +
  labs(title= "Correlation Between Transaction Count and Median of Days Since Prior Order") +
  xlab("Transaction Count") +
  ylab("Median of Days Since Prior Order")
```



```
## `geom_smooth()` using method = 'gam'
```



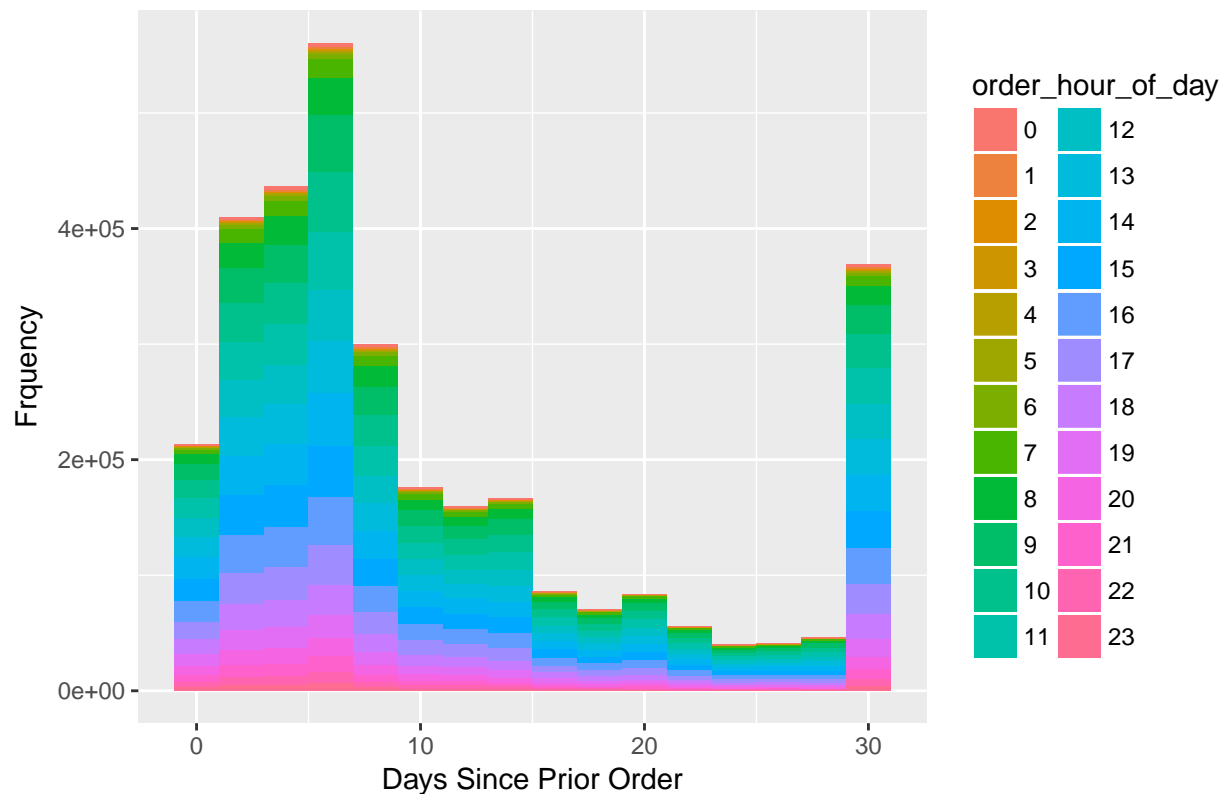
We clearly see that customers who have more transactions tend to shop more.

## Distribution of Days Since Prior Order

```
ggplot(orders, aes(days_since_prior_order, fill = order_hour_of_day)) +  
  geom_histogram(binwidth = 2) +  
  labs(title = "Distribution of Days Since Prior Order, Colored by Order Hour of Day") +  
  xlab("Days Since Prior Order") +  
  ylab("Frquency") +  
  theme(plot.title = element_text(hjust = 0.5))
```

```
## Warning: Removed 206209 rows containing non-finite values (stat_bin).
```

## Distribution of Days Since Prior Order, Colored by Order Hour of Day



```
#transaction made monthly apart
prop.table(table(orders$days_since_prior_order == 30))
```

```
##
##      FALSE      TRUE
## 0.8851205 0.1148795
```

```
#transaction made within interval 0 days to 10 days
prop.table(table(orders$days_since_prior_order > 0 & orders$days_since_prior_order <= 10))
```

```
##
##      FALSE      TRUE
## 0.3942627 0.6057373
```

From the plot above, we see a bimodal distribution. We can see that there are around 380K (11.48%) transactions made 30 days (one month) after the prior transaction. Around 63% of the customers make purchase ranging from 0 day to 10 days (around 2% of customer make repurchase on the same day).


## When customers come back, do they make purchase around the same time (day/hour) ?

We also wonder when customers come back to make another purchase, is there a consistent trend? For example, a customer made purchase this Monday at 10AM, will this time frame be somewhat similar the next time this customer come back to make a purchase again ? We will find out in this section.

We will use standard deviation as a metric to gauge the consistency of specific customer.

```
#Calculate standard deviation to both of the "order_dow" and "order_hour_of_day", based on user_id, pas
order_patterns <- orders %>% select(user_id, order_dow, order_hour_of_day) %>%
  group_by(user_id) %>%
  summarise(dow_sd = sd(order_dow), order_hour_of_day_sd = sd(as.numeric(order_hour_of_d

boxplot(order_patterns$dow_sd)
```



```
quantile(order_patterns$dow_sd)
```


```
##      0%      25%      50%      75%     100%
## 0.000000 1.576482 1.949359 2.267787 3.464102
```

50% of the customer make purchase 1.58 to 2.27 standard deviation of their day of the week means. We have two hypotheses to explain this phenomenon:

1. Only a small percentage of customer make purchase on a specific day of the week, most customers would just make purchase whenever they need, and products they purchase are different every time (reordering is low).
2. Only a small percentage of customer subscribe to a certain type of products and scheduled delivery for certain interval of time, most customers make purchase when products they often buy are exhausted, however, because most people would wait a day or two before placing new orders, which explains the standard deviation.

We will try to address this observation in later section, when we combine the product information and the order history into one data frame.

```
boxplot(order_patterns$order_hour_of_day_sd)
```



```
quantile(order_patterns$order_hour_of_day_sd)
```

```
##           0%          25%          50%          75%         100%
##  0.000000  2.629956  3.473342  4.440077 12.996794
```

Comparing to the day of the week standard deviations, the standard deviation for hour of the day is higher, 50% of the customers make purchase between 2.63 to 4.44 standard deviations of their means in hour of the day. This huge standard deviation make sense since we don't expect customers would keep track when they place order last time, and try to place the next order around the same time.

## Orders\_products (prior and train)

There are different data files containing the order history in every trasaction, file “order\_products\_prior” and file “order\_products\_train”. These two files contain variable “order\_id”, “product\_id”, “add\_to\_cart\_order”, and a binary variable “reordered”, with 1 indicates yes and 0 indicates no. In order to get a full grasp of all the transactions, we will combine these two data at the moment.

```
#combining of the prior and train files
```

```
combined <- rbind(order_products_prior, order_products_train)
dim(combined)
```

```
## [1] 33819106      4
```

```
str(combined)
```

```
## Classes 'data.table' and 'data.frame': 33819106 obs. of 4 variables:
```

```
## $ order_id : int 2 2 2 2 2 2 2 2 2 3 ...
```

```
## $ product_id : int 33120 28985 9327 45918 30035 17794 40141 1819 43668 33754 ...
```

```
## $ add_to_cart_order: int 1 2 3 4 5 6 7 8 9 1 ...
```

```
## $ reordered : int 1 1 0 1 0 1 1 1 0 1 ...
```

```
## - attr(*, ".internal.selfref")=<externalptr>
```

We will then use the product ID to get the product names, and store them under variable “product”.

```
combined$product <- products$product_name[combined$product_id]
combined
```

```
##      order_id product_id add_to_cart_order reordered
##      1:      2      33120                1         1
```

```
##      2:      2      28985      2      1
##      3:      2      9327      3      0
##      4:      2     45918      4      1
##      5:      2     30035      5      0
##      ---
## 33819102: 3421063     14233      3      1
## 33819103: 3421063     35548      4      1
## 33819104: 3421070     35951      1      1
## 33819105: 3421070     16953      2      1
## 33819106: 3421070      4724      3      1
##
##                                product
##      1:      Organic Egg Whites
##      2:      Michigan Organic Kale
##      3:      Garlic Powder
##      4:      Coconut Butter
##      5:      Natural Sweetener
##      ---
## 33819102:      Natural Artesian Water
## 33819103:      Twice Baked Potatoes
## 33819104: Organic Unsweetened Almond Milk
## 33819105:      Creamy Peanut Butter
## 33819106:      Broccoli Florettes
```

## Most popular products

We will then sort the the product and see which products are the most popular.

```
product_count <- data.frame(product_name = table(combined$product) %>% sort(decreasing = T) %>%
                             names,
                             count = table(combined$product) %>% sort(decreasing = T) %>% unlist
                             %>% as.numeric())
```

Here are the ten most popular products in the grocery:

```
product_count %>% head(10)

##      product_name  count
## 1      Banana 491291
## 2 Bag of Organic Bananas 394930
## 3  Organic Strawberries 275577
## 4  Organic Baby Spinach 251705
## 5  Organic Hass Avocado 220877
## 6    Organic Avocado 184224
## 7      Large Lemon 160792
## 8      Strawberries 149445
## 9          Limes 146660
## 10 Organic Whole Milk 142813
```

People really love vegetables, fruits, and organic products.

## First product add to the cart

Using the variable “add\_to\_cart\_order”, we can find out which items usually get to added to the cart first.

```
combined[add_to_cart_order == 1, .N, by = product][order(-N)] %>% head(10)
```

```
##           product      N
## 1:           Banana 115521
## 2: Bag of Organic Bananas 82877
## 3:      Organic Whole Milk 32071
## 4:      Organic Strawberries 28875
## 5:      Organic Hass Avocado 24913
## 6:      Organic Baby Spinach 24412
## 7:           Organic Avocado 23393
## 8:           Spring Water 17552
## 9:           Strawberries 17073
## 10:      Organic Raspberries 14950
```

We can see that bananas are still the winner.

## Which products get reordered most

```
combined %>%
  select(product, reordered) %>%
  group_by(product) %>%
  summarise(reorder_sum = sum(reordered)) %>%
  arrange(desc(reorder_sum)) %>%
  head(10)
```

```
## # A tibble: 10 x 2
##           product reorder_sum
##           <fctr>      <int>
## 1           Banana      415166
## 2 Bag of Organic Bananas 329275
## 3      Organic Strawberries 214448
## 4      Organic Baby Spinach 194939
## 5      Organic Hass Avocado 176173
## 6           Organic Avocado 140270
## 7      Organic Whole Milk 118684
## 8           Large Lemon   112178
## 9      Organic Raspberries 109688
## 10          Strawberries 104588
```

This top 10 list is pretty similar to the most popular products, which make sense, given the fact that nearly 60% of the products are from reordered transactions.

## In-depth Analysis

### Combining everything

From this point forward, we will try to combined all the information we have at this point and try to gain a deeper insight of the transaction history.

```
combined$order_id <- as.factor(combined$order_id)
```

```
order_hist_detail <- merge(combined, orders, by = "order_id", all.x = TRUE)
#remove repeated variable
```

```
order_hist_detail$product <- NULL
```

```
order_hist_detail %>% head(10)
```

```
##      order_id product_id add_to_cart_order reordered user_id eval_set
## 1:         1      49302             1         1 112108    train
## 2:         1      11109             2         1 112108    train
## 3:         1      10246             3         0 112108    train
## 4:         1      49683             4         0 112108    train
## 5:         1      43633             5         1 112108    train
## 6:         1      13176             6         0 112108    train
## 7:         1      47209             7         0 112108    train
## 8:         1      22035             8         1 112108    train
## 9:         2      33120             1         1 202279    prior
## 10:        2      28985             2         1 202279    prior
##      order_number order_dow order_hour_of_day days_since_prior_order
## 1:              4         4             10              9
## 2:              4         4             10              9
## 3:              4         4             10              9
## 4:              4         4             10              9
## 5:              4         4             10              9
## 6:              4         4             10              9
## 7:              4         4             10              9
## 8:              4         4             10              9
## 9:              3         5              9              8
## 10:             3         5              9              8
```

Combining transaction history with product information. Since we have almost fifty thousand products, adding the aisle and department information can help use to narrow down customer's preferences.

```
order_hist_detail <- merge(order_hist_detail, products_w_desc, by = "product_id")[order(order_id)]
order_hist_detail %>% head(10)
```

```
##      product_id order_id add_to_cart_order reordered user_id eval_set
## 1:      10246         1             3         0 112108    train
## 2:      11109         1             2         1 112108    train
## 3:      13176         1             6         0 112108    train
## 4:      22035         1             8         1 112108    train
## 5:      43633         1             5         1 112108    train
## 6:      47209         1             7         0 112108    train
## 7:      49302         1             1         1 112108    train
## 8:      49683         1             4         0 112108    train
## 9:       1819         2             8         1 202279    prior
## 10:       9327         2             3         0 202279    prior
##      order_number order_dow order_hour_of_day days_since_prior_order
## 1:              4         4             10              9
## 2:              4         4             10              9
## 3:              4         4             10              9
## 4:              4         4             10              9
## 5:              4         4             10              9
## 6:              4         4             10              9
## 7:              4         4             10              9
## 8:              4         4             10              9
## 9:              3         5              9              8
## 10:             3         5              9              8
```

```
##           product_name aisle_id department_id
## 1:           Organic Celery Hearts      83         4
## 2: Organic 4% Milk Fat Whole Milk Cottage Cheese    108        16
## 3:           Bag of Organic Bananas      24         4
## 4:           Organic Whole String Cheese      21        16
## 5:           Lightly Smoked Sardines in Olive Oil    95        15
## 6:           Organic Hass Avocado      24         4
## 7:           Bulgarian Yogurt     120        16
## 8:           Cucumber Kirby      83         4
## 9:           All Natural No Stir Creamy Almond Butter    88        13
## 10:           Garlic Powder     104        13
##      aisles_description depart_description
## 1:      fresh vegetables      produce
## 2: other creams cheeses      dairy eggs
## 3:      fresh fruits      produce
## 4:      packaged cheese      dairy eggs
## 5:  canned meat seafood  canned goods
## 6:      fresh fruits      produce
## 7:           yogurt      dairy eggs
## 8:      fresh vegetables      produce
## 9:           spreads      pantry
## 10:      spices seasonings      pantry
```

## Most popular products in each department

There are twenty departments total (excluding “missing” department), we wonder what are the most popular (top three) products in each department, and how do their transaction count fluctuate as a day progress.

```
top_3_each_dept = list()
for (i in 1:(nrow(dept)-1)) {
  tmp_df <- order_hist_detail[depart_description == dept$department[i]] %>%
    select(aisles_description) %>%
    group_by(aisles_description) %>%
    summarise(count = n()) %>%
    arrange(desc(count))

  top_3_each_dept[[i]] = tmp_df$aisles_description[1:3]
}
```

```
tmp = top_3_each_dept %>% unlist
tmp = tmp[!is.na(tmp)]
tmp
```

```
## [1] frozen produce      ice cream ice
## [3] frozen meals        other
## [5] bread               breakfast bakery
## [7] tortillas flat bread fresh fruits
## [9] fresh vegetables    packaged vegetables fruits
## [11] beers coolers       red wines
## [13] white wines         asian foods
## [15] latino foods        indian foods
## [17] water seltzer sparkling water refrigerated
## [19] soft drinks         cat food care
## [21] dog food care       dry pasta
```



```

## [23] pasta sauce                instant foods
## [25] bulk dried fruits vegetables bulk grains rice dried goods
## [27] soap                      oral hygiene
## [29] vitamins supplements      hot dogs bacon sausage
## [31] poultry counter           packaged poultry
## [33] baking ingredients        spreads
## [35] oils vinegars             cereal
## [37] hot cereal pancake mixes  granola
## [39] soup broth bouillon       canned jarred vegetables
## [41] canned meals beans        yogurt
## [43] packaged cheese           milk
## [45] paper goods               cleaning products
## [47] laundry                   baby food formula
## [49] diapers wipes             baby bath body care
## [51] chips pretzels            crackers
## [53] energy granola bars       lunch meat
## [55] fresh dips tapenades      tofu meat alternatives
## 134 Levels: air fresheners candles asian foods ... yogurt

index <- which(order_hist_detail$aisles_description %in% tmp) %>% unlist %>% as.numeric()
tmp2 <- order_hist_detail[index,]


hourly_total = tmp2[, .(hourly_total = .N), by = order_hour_of_day]

tmp3 <- tmp2[, .N, by = .(aisles_description, depart_description, order_hour_of_day)]
tmp3 <- merge(tmp3, hourly_total, by = "order_hour_of_day")
tmp3 <- tmp3 %>% mutate(pct = N/hourly_total*100)

tmp3$order_hour_of_day <- as.numeric(tmp3$order_hour_of_day)

ggplot(tmp3, aes(order_hour_of_day, pct, group = interaction(aisles_description, depart_description))) +
  geom_area(color = "black", aes(fill = depart_description))

```



From the plot above we can see how the percentage of sales in different aisels move as the day progress. Generally speaking, we don't see the percentage of sales shift too much. In addition, we also see that aisles produce, snacks take up approximately 50% of the sales, among all the sales of top 3 products in each department.

```
tmp <- products_w_desc
depart_size <- tmp[, .(depart_size = .N), by = depart_description]
depart_aisles_size <- tmp[, .(aisles_size = .N), by = .(aisles_description, depart_description)]



depart_aisles_size <- merge(depart_aisles_size, depart_size, by = "depart_description", all.x = TRUE)

item_sales_count <- order_hist_detail[, .(item_sale_count = .N), by = aisles_description]

treemap_df <- merge(depart_aisles_size, item_sales_count, by = "aisles_description")

treemap(treemap_df,
  index = c("depart_description", "aisles_description"),
  vSize = "aisles_size",
  vColor = "item_sale_count",
  type = "value",
  title = "Aisle and Product Treemap, Sized by Aisle Size, Colored by Sale Count",
  title.legend = "Sales Count")
```

Aisle and Product Treemap, Sized by Aisle Size, Colored by Sale Count



Earlier we've made an assumption that because "personal care" and "snacks" are the department carry most products in the grocery store, they should have the highest sale, however, we were wrong. From the treemap, we see that department that sold most products were the produce department (thanks to those bananas).